

pointcloudlibrary

PCL :: Search

Marius Muja and Julius Kammerl

July 1, 2011


- ▶ Nearest neighbor search is an inner loop of many parts of PCL (filters, surface, features, registration)
 - ▶ Needs to be as fast as possible
- ▶ **FLANN** - Fast Library for Approximate Nearest Neighbors
 - ▶ <http://www.cs.ubc.ca/~mariusm/flann>
 - ▶ C, C++, Matlab and Python bindings
 - ▶ Exact nearest neighbor search in low dimensional spaces (3D) using kd-trees
 - ▶ Approximate nearest neighbor search in high dimensional spaces
- ▶ Octree - 3D search





- ▶ Nearest neighbor search problem
 - ▶ Given a set of points $P = p_1, p_2, \dots, p_n$ in a metric space X , preprocess them in such a way that given a new point $q \in X$ finding the closest p_i to q can be done easily
- ▶ K-Nearest neighbor search
 - ▶ find the closest K neighbors
- ▶ Radius nearest neighbor search
 - ▶ find all the neighbors within a certain radius


1. KdTree
2. 3D Nearest Neighbor Search
3. High Dimensional Nearest Neighbor Search
4. Octree


- ▶ recursively divide the data points based on a single dimension
 - ▶ how to choose the dimension in which to divide the data?
 - ▶ where to divide?
- ▶ binary tree
- ▶ when searching entire branches can be ignored due to being too far away from the query point
- ▶ very efficient for low dimensionality data




















- ▶ during 2011 GSOC, 2 students are working on a GPU based kd-tree implementation
- ▶ encouraging preliminary results, speedups of 8-10x compared to CPU implementation



- ▶ Header: `#include <pcl/kdtree/kdtree_flann.h>`
- ▶ Class: `template<typename PointT> class pcl::KdTreeFLANN`
- ▶ K-nearest neighbor search

```
int nearestKSearch (const PointT &point, int k,
                    vector<int> &k_indices, vector<float> &k_distances);
```
- ▶ Radius search

```
int radiusSearch (const PointT &point, double radius,
                  vector<int> &k_indices,
                  vector<float> &k_distances, int max_nn = -1);
```

```
PointCloud<PointXYZ>::Ptr cloud (new PointCloud<PointXYZ>);
PointXYZ searchPoint;

// ... populate the cloud and the search point

// create a kd-tree instance
KdTreeFLANN<PointXYZ> kdtree;

// assign a point cloud - this builds the tree
kdtree.setInputCloud (cloud);

// pre-allocate the neighbor index and
// distance vectors
int K = 10;
std::vector<int> pointsIdx(K);
std::vector<float> pointsSquaredDist(K);

// K nearest neighbor search
kdtree.nearestKSearch (searchPoint, K, pointsIdx, pointsSquaredDist);
```




```
PointCloud<PointXYZ>::Ptr cloud (new PointCloud<PointXYZ>);
PointXYZ searchPoint;

// ... populate the cloud and the search point

// create a kd-tree instance
KdTreeFLANN<PointXYZ> kdtree;

// assign a point cloud - this builds the tree
kdtree.setInputCloud (cloud);

std::vector<int> pointIdxRadius;
std::vector<float> pointsSquaredDistRadius;
float radius = ...;

// radius search
int count = kdtree.radiusSearch (searchPoint, radius,
                                pointIdxRadiusSearch, pointsSquaredDistRadius);
```



Compile & Try

```
$ cd $PCL_ROOT/doc/tutorials/content/sources/kdtree_search
$ mkdir build
$ cd build
$ cmake ..
$ make
$ ./kdtree_search
K nearest neighbor search at (701.248 662.202 554.841) with K=10
  702.91 601.583 521.043 (squared distance: 4819.7)
  676.792 699.92 482.203 (squared distance: 7297.07)
  731.215 717.665 491.714 (squared distance: 7959.15)
  670.142 707.355 476.051 (squared distance: 9214.31)
  681.636 728.872 479.31 (squared distance: 10534.5)
  683.843 581.742 492.494 (squared distance: 10663.9)
  696.085 705.888 457.71 (squared distance: 11369.7)
  683.603 667.109 430.477 (squared distance: 15801.9)
  721.228 684.503 430.334 (squared distance: 16398.5)
  829.566 676.396 560.64 (squared distance: 16700.6)
Neighbors within radius search at (701.248 662.202 554.841) with
radius=114.069
  702.91 601.583 521.043 (squared distance: 4819.7)
  676.792 699.92 482.203 (squared distance: 7297.07)
  731.215 717.665 491.714 (squared distance: 7959.15)
  670.142 707.355 476.051 (squared distance: 9214.31)
  681.636 728.872 479.31 (squared distance: 10534.5)
  683.843 581.742 492.494 (squared distance: 10663.9)
  696.085 705.888 457.71 (squared distance: 11369.7)
```

1. KdTree
2. 3D Nearest Neighbor Search
3. High Dimensional Nearest Neighbor Search
4. Octree

► Object recognition (TOD)




- ▶ Object recognition (TOD)
- ▶ Image stitching (AutoStitch, Hugin)



AutoStitch ©Matthew Brown

- ▶ Object recognition (TOD)
- ▶ Image stitching (AutoStitch, Hugin)
- ▶ 3D Reconstruction (Photosynth)



- ▶ Object recognition (TOD)
- ▶ Image stitching (AutoStitch, Hugin)
- ▶ 3D Reconstruction (Photosynth)
- ▶ 3D object classification (VFH)



VFH Classification example



- ▶ Object recognition (TOD)
- ▶ Image stitching (AutoStitch, Hugin)
- ▶ 3D Reconstruction (Photosynth)
- ▶ 3D object classification (VFH)
- ▶ Content based image retrieval

- ▶ Object recognition (TOD)
- ▶ Image stitching (AutoStitch, Hugin)
- ▶ 3D Reconstruction (Photosynth)
- ▶ 3D object classification (VFH)
- ▶ Content based image retrieval
- ▶ Visual SLAM




- ▶ For high dimensionality data, no exact algorithm faster than linear search is known
- ▶ Approximate nearest neighbor search is used to obtain large speedups
- ▶ FLANN contains several algorithms for high dimensional approximate nearest neighbor search
 - ▶ KDTreeIndex (randomized kd-tree forest)
 - ▶ KMeansIndex (hierarchical k-means tree)
 - ▶ HierarchicalClusteringIndex (clustering tree in a generic metric space)*
 - ▶ LshIndex (locality sensitive hashing)*

- ▶ multiple trees are build in parallel
- ▶ the split dimension is chosen randomly from the first D dimensions with greatest variance ($D=5$)
- ▶ at search time a single priority queue is used across all trees
- ▶ search is terminated after a predefined number of tree leafs are checked



- ▶ Building the tree
 - ▶ built by splitting the data at each level of the tree using k-means clustering
 - ▶ apply the same procedure recursively on each cluster
 - ▶ just a few iterations of the k-means clustering give good results
- ▶ Exploring the tree
 - ▶ unexplored branches are added to a priority queue while traversing the tree
 - ▶ restart search from best branch in the priority queue



(Nistér & Stewénius, 2006)

- ▶ Header: `#include <flann/flann.h>`
- ▶ Class: `template<typename Distance> class flann::Index`
- ▶ K-nearest neighbor search

```
void knnSearch(const Matrix<ElementType>& queries,  
               Matrix<int>& indices, Matrix<DistanceType>& dists,  
               int knn, const SearchParams& params)
```

- ▶ Radius search

```
int radiusSearch(const Matrix<ElementType>& query,  
                 Matrix<int>& indices, Matrix<DistanceType>& dists,  
                 float radius, const SearchParams& params)
```

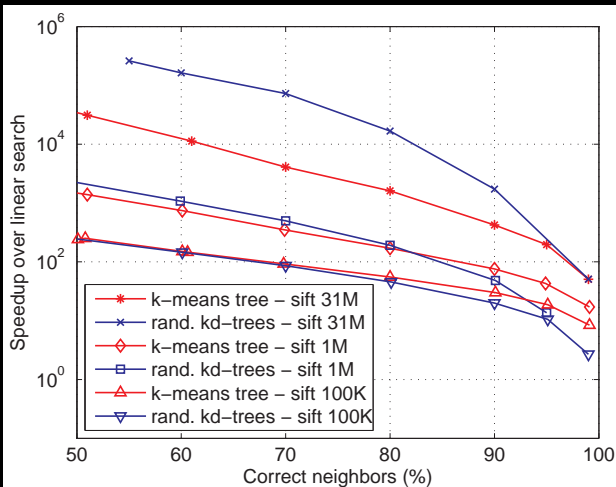
```
flann::Matrix<float> data;
flann::Matrix<float> queries;
// populate the matrix with features
// one feature/row

// build randomized kd-tree index (4 trees)
flann::Index< L2<float> > index(data, flann::KdTreeIndexParams(4));
index.buildIndex();

// allocate memory for results
int k = 10;
int n = queries.rows;
flann::Matrix<int> k_indices(new int[n*k], n, k);
flann::Matrix<float> k_distances(new float[n*k], n, k);

// KNN search
index.knnSearch (queries, k_indices, k_distances, k,
                 flann::SearchParams(256));
```

- Speedup with precision for different dataset sizes








- ▶ See tutorial at: http://www.pointclouds.org/documentation/tutorials/vfh_recognition.php


```
$ cd $PCL_ROOT/doc/tutorials/content/sources/vfh_recognition
$ wget http://dev.pointclouds.org/attachments/download/216/
vfh_recognition_tutorial_data.tbz
$ tar -xzf vfh_recognition_tutorial_data.tbz
$ mkdir build
$ cd build
$ cmake ..
$ make
$ cd ..
$ ./build/build_tree data
...
$ ./build/nearest_neighbors -k 16 -thresh 50 data/000.580.67/
1258730231333_cluster_0_nxyz_vfh.pcd
```

Octree - 3D hierarchical spatial tree data structure

- ▶ Recursive divide & conquer algorithm
- ▶ Binary subdivision of occupied cells into 8 octants (voxels)

2D Example (Quadtree):





- ▶ Root node describes a cubic bounding box which encapsulates all points
- ▶ Child nodes recursively subdivide point space
- ▶ Nodes have up to eight children \Rightarrow Byte encoding

Instantiate octree:

```
float voxelSize = 0.01f; // voxel resolution
OctreePointCloud<PointXYZ> octree (voxelSize);
```

Set input point cloud (via Boost shared pointers):

```
octree.setInputCloud (cloud);
```

Define octree bounding box (optional):

```
// calculate bounding box of input cloud
octree.defineBoundingBox ();
// manually define bounding box
octree.defineBoundingBox (minX, minY, minZ, maxX, maxY, maxZ);
```

Add points from input cloud to octree:

```
octree.addPointsFromInputCloud ();
```

Delete octree data structure:

(pushes allocated nodes to memory pool!)

```
octree.deleteTree ();
```

Check if voxel at given point coordinates exist:

```
double X,Y,Z;  
bool occupied;  
X = 1.0; Y=2.0; Z=3.0;  
occupied = octree.isVoxelOccupiedAtPoint (X, Y, Z);
```

Get center points of all occupied voxels:
(voxel grid filter/downsampling)

```
std::vector<PointXYZ> pointGrid;  
octree.getOccupiedVoxelCenters (pointGrid);
```

Delete voxel:


```
pcl::PointXYZ point_arg( 1.0, 2.0, 3.0 );  
octree.deleteVoxelAtPoint ( point );
```

Provided algorithms in PCL using octrees for spatial decomposition:

- ▶ Search operations (neighbor search, radius search, voxel search)
- ▶ Downsampling (Voxel-grid / Voxel-centroid filter)
- ▶ Point cloud compression
- ▶ Spatial change detection
- ▶ Spatial point density analysis
- ▶ Occupancy checks/maps
- ▶ Collision detection
- ▶ ...

Points within radius search

- ▶ Depth first tree exploration
- ▶ At every node investigate occupied child voxels that overlap with search sphere



K nearest neighbor search:

- ▶ Priority queue (binary heap) of nodes and point candidates
- ▶ Investigate occupied child voxels (closest voxel first)
- ▶ Radius search with radius=distance to Kth point candidate
- ▶ Update radius with every new point candidate

Define search precision / error bound:

```
octree.setEpsilon (double eps); // default: 0.0
```

Neighbors within voxel search:

```
std::vector<int> pointIdxVec;  
  
if (octree.voxelSearch (searchPoint, pointIdxVec))  
{  
    for (size_t i = 0; i < pointIdxVec.size (); ++i)  
        std::cerr << " " << cloud->points[pointIdxVec[i]].x  
            << " " << cloud->points[pointIdxVec[i]].y  
            << " " << cloud->points[pointIdxVec[i]].z << std::endl;  
}
```

K nearest neighbor search:

```
int K = 10;  
std::vector<int> pointIdxNKNSearch;  
std::vector<float> pointNKNSquaredDistance;  
  
if ( octree.nearestKSearch (searchPoint, K,  
    pointIdxNKNSearch, pointNKNSquaredDistance) > 0 )  
{  
    ...  
}
```

Neighbors within radius search:


```
std::vector<int> pointIdxRadiusSearch;  
std::vector<float> pointRadiusSquaredDistance;  
float radius = 0.1;  
  
if ( octree.radiusSearch (searchPoint, radius,  
    pointIdxRadiusSearch, pointRadiusSquaredDistance) > 0 )  
{  
    ...  
}
```

Approx. neighbors within radius search:

(only scans points within “search point voxel”)

```
std::vector<int> pointIdxRadiusSearch;  
std::vector<float> pointRadiusSquaredDistance;  
float radius = 0.1;  
  
if ( octree.approxNearestSearch (searchPoint, radius,  
    pointIdxRadiusSearch, pointRadiusSquaredDistance) > 0 )  
{  
    ...  
}
```

Template configuration:



Optimized performance&memory usage:

- ▶ Select octree base implementation
- ▶ Select/define leaf node class
- ▶ Serialization callbacks (serializeLeafCallback, deserializeLeafCallback, serializeNewLeafCallback)

OctreePointCloud classes:

```
float resolution = 0.01f;  
  
// equal to OctreePointCloudPointVector<PointXYZ>  
OctreePointCloud<PointXYZ> octreeA (resolution);  
  
// manages indices vectors in leaf nodes  
OctreePointCloudPointVector<PointXYZ> octreeB (resolution);  
// keeps a single point indices in leaf nodes  
OctreePointCloudSinglePoint<PointXYZ> octreeC (resolution);  
// does not store any point information in leaf node  
OctreePointCloudOccupancy<PointXYZ> octreeD (resolution);
```

Octree-Base selection via typedefs:

```
OctreePointCloud<PointXYZ>::SingleBuffer octreeSB (resolution);  
OctreePointCloud<PointXYZ>::DoubleBuffer octreeDB (resolution);  
OctreePointCloud<PointXYZ>::LowMem octreeLM (resolution);
```

Octree2BufBase implementation:

- ▶ Create octrees at high rate
- ▶ Advanced memory management:
 - ▶ Previous tree structure is kept in memory
 - ▶ Maximum reuse of already allocated branch&leaf nodes
 - ▶ Unused node instances are pushed to a memory pool for later reuse
- ▶ Enables comparison of octree structure (change detection)

Switching between octree buffers:

```
octree.switchBuffers ();
```




Change Detection

```
class SimpleSpatialChangeDetection
{
public:
    OctreePointCloudChangeDetector<PointXYZRGB>* octree;
    ...
    void
    cloud_cb_ (const pcl::PointCloud<pcl::PointXYZRGB>::ConstPtr &cloud)
    {
        if (!viewer.wasStopped ())
        {
            // Switch octree buffers
            octree.switchBuffers ();

            // Add points from cloud to octree
            octree.setInputCloud (cloud);
            octree.addPointsFromInputCloud ();

            std::vector<int> newPointIdxVector;

            /* Get vector of point indices from octree voxels
               which did not exist in previous buffer */
            octree.getPointIndicesFromNewVoxels (newPointIdxVector);
            ...
        }
    }
};
```



- ▶ Real-time spatial change detection based on XOR comparison of octree structure

DEMO: See /visualization/tool/openni_change_viewer

Example: Point density estimation

Design your own leaf node class:

```
template<typename DataT>
class OctreePointCloudDensityLeaf : public OctreeLeafAbstract<DataT>
{
public:
    ...

    virtual void
    setData (const DataT& point_arg)
    {
        pointCounter_++;
    }

    unsigned int
    getPointCounter ()
    {
        return pointCounter_;
    }

    ...

private:
    unsigned int pointCounter_;
};
```


.. and your own OctreePointCloud class:

```
class OctreePointCloudDensity : public OctreePointCloud
<PointT, OctreePointCloudDensityLeaf<int> , OctreeT>
{
public:
    ...


    unsigned int
    getVoxelDensityAtPoint (const PointT& point_arg) const
    {
        unsigned int pointCount = 0;

        OctreePointCloudDensityLeaf<int>* leaf =
            this->findLeafAtPoint (point_arg);

        if (leaf) pointCount = leaf->getPointCounter ();

        return pointCount;
    }
};
```

Encoding Pipeline:



Example:

```

/* for a full list of profiles see:
   /io/include/pcl/compression/compression_profiles.h */
compression_Profiles_e compressionProfile =
  pcl::octree::MED_RES_ONLINE_COMPRESSION_WITH_COLOR;

// instantiate point cloud compression for encoding and decoding
PointCloudCompression<PointXYZ> PointCloudEncoder (compressionProfile)
PointCloudCompression<PointXYZ> PointCloudDecoder ();
...
// iostream to read/write compressed point cloud data
std::stringstream compressedData;

// compress & decompress point cloud
PointCloudEncoder->encodePointCloud (cloud, compressedData);
PointCloudDecoder->decodePointCloud (compressedData, cloudOut);

```

- ▶ See octree search tutorial at:
`http://pointclouds.org/documentation/tutorials/octree.php`
- ▶ See point cloud compression tutorial at:
`http://pointclouds.org/documentation/tutorials/compression.php`
- ▶ See change detection tutorial at:
`http://pointclouds.org/documentation/tutorials/octree_change.php`
- ▶ Point cloud compression and streaming app:
`PCL_ROOT/apps/openni_stream_compression`
- ▶ Change detection app:
`PCL_ROOT/visualization/tools/openni_change_viewer`